



OrganSlot E: Cross-Attention



Offline joint training | 20 tokens per organ | 8 organs + background slot

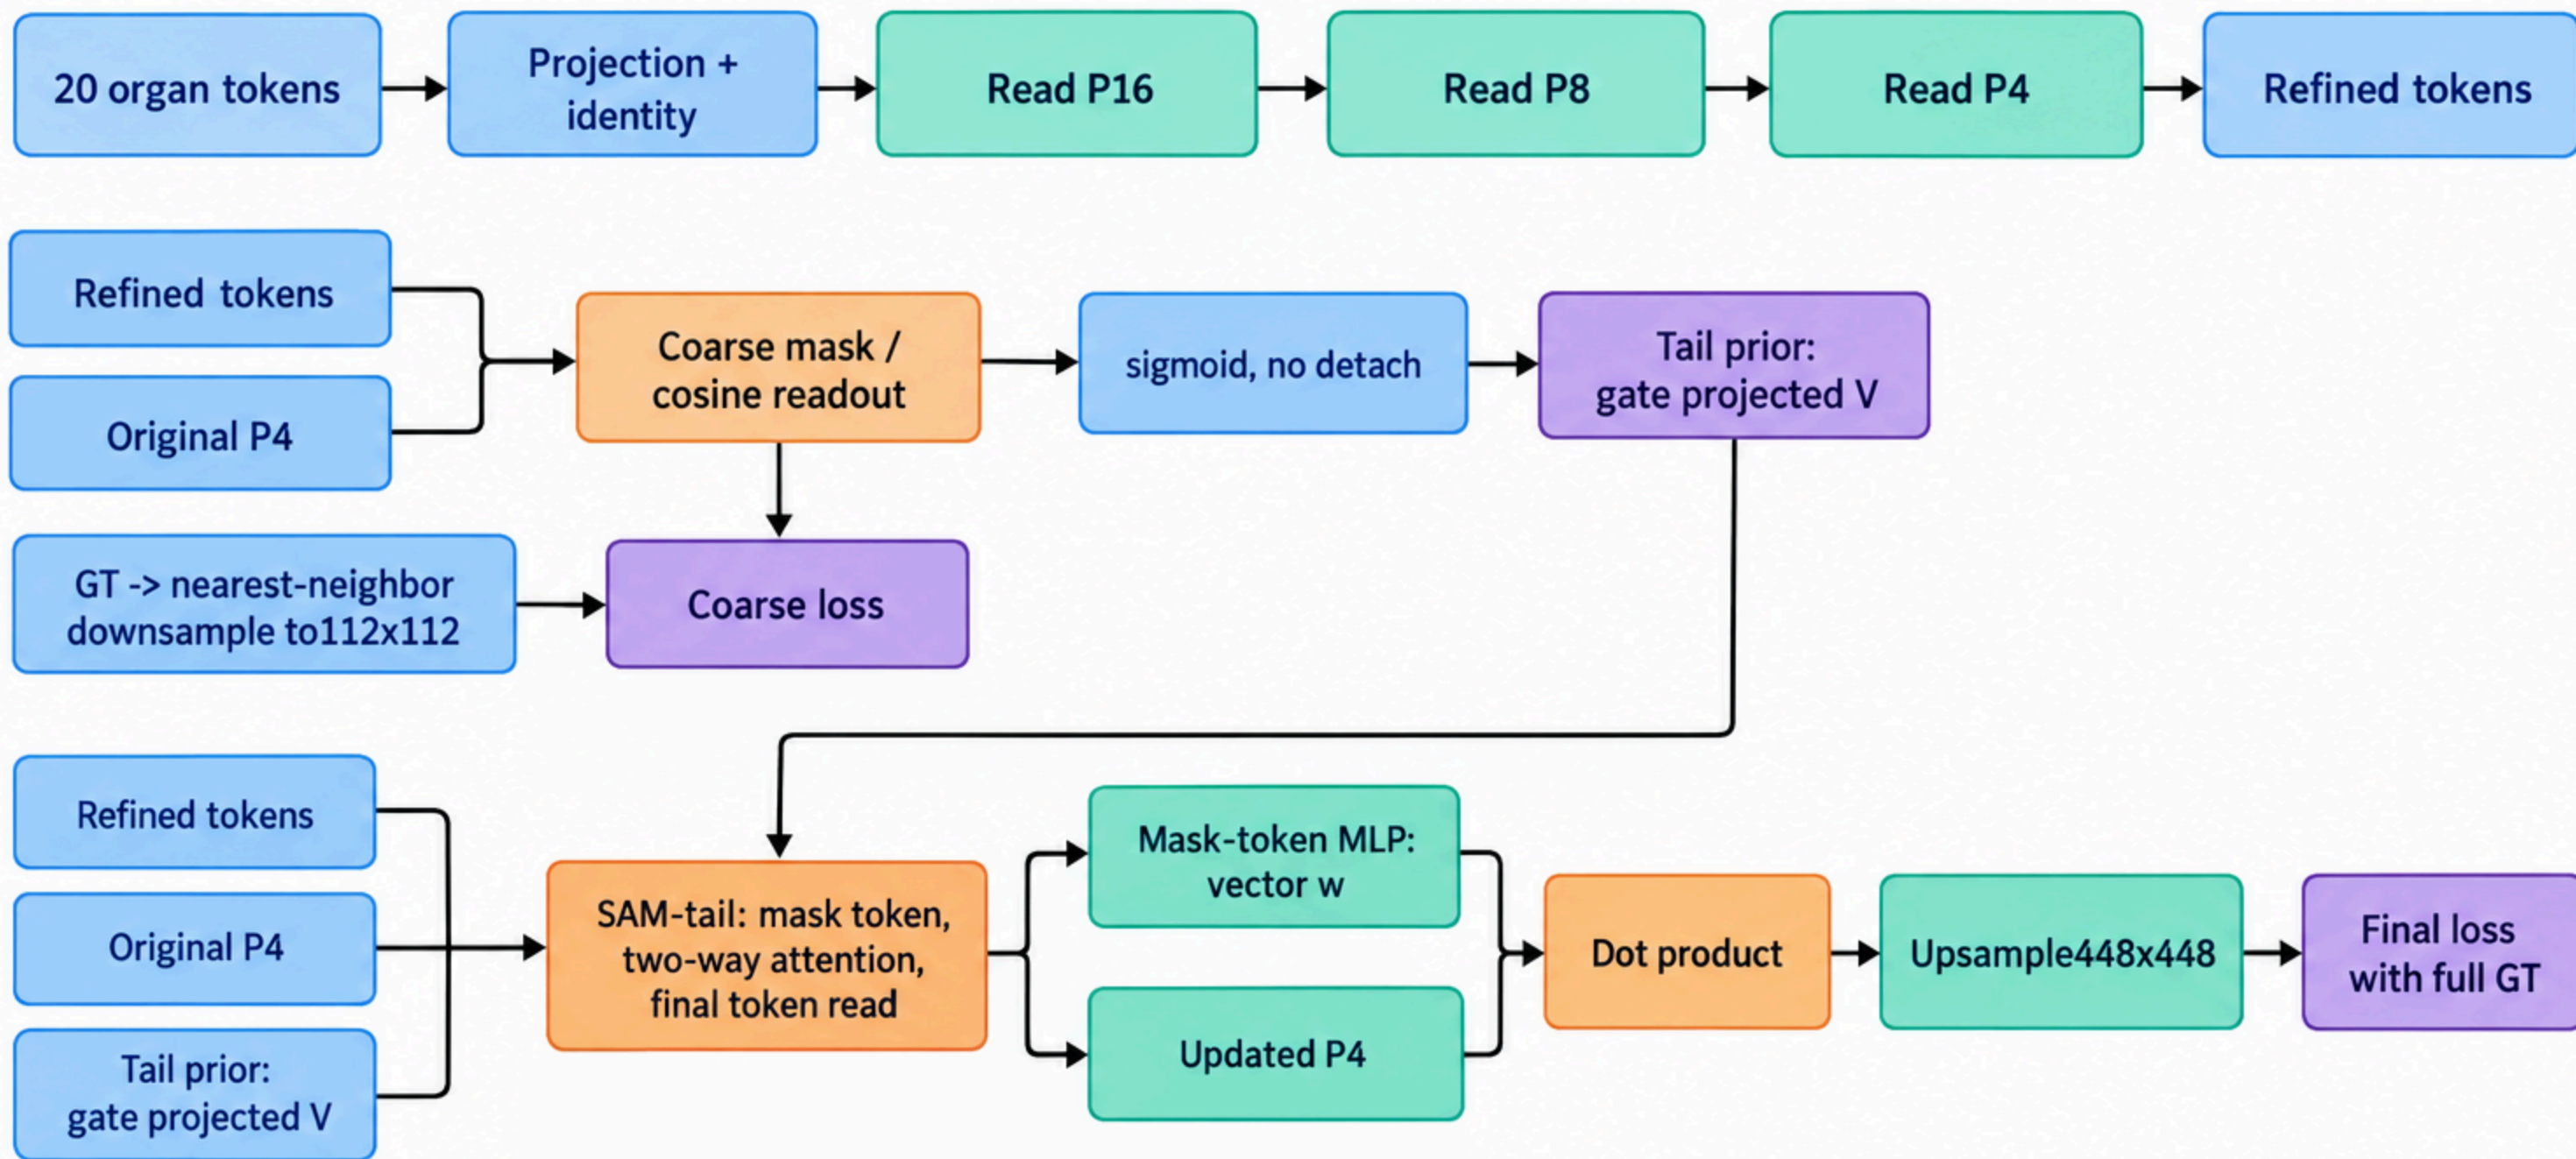
SAM-tail: Soft Prior + Aux



No coarse-mask prior inside **SAM-tail**

$L_{seg} = 0.01 (L_{final} + 0.25 \text{ mean } L_{aux})$

SAM-tail: Downsampled-GT Coarse-to-Fine



Spatial source: CT -> shared ViT and shallow CNN -> shared pixel decoder -> P16, P8, P4.

Token source: ViT z -> OrganCollector -> TGEnc -> 20 tokens.

Reconstruction: original tokens -> AHER -> canvas fusion -> shared reconstruction decoder.

Two decoder stages in one forward pass, not two training stages.